



Services :

Each pod gets its own IP address however in a deployment the set of pods running in a moment of time could be different from the set of pods running the application the moment later



If lets the backend pod get terminated ? The new pod would created by replication controller ? yes

But when a new pod is created it wil have a different ip address.



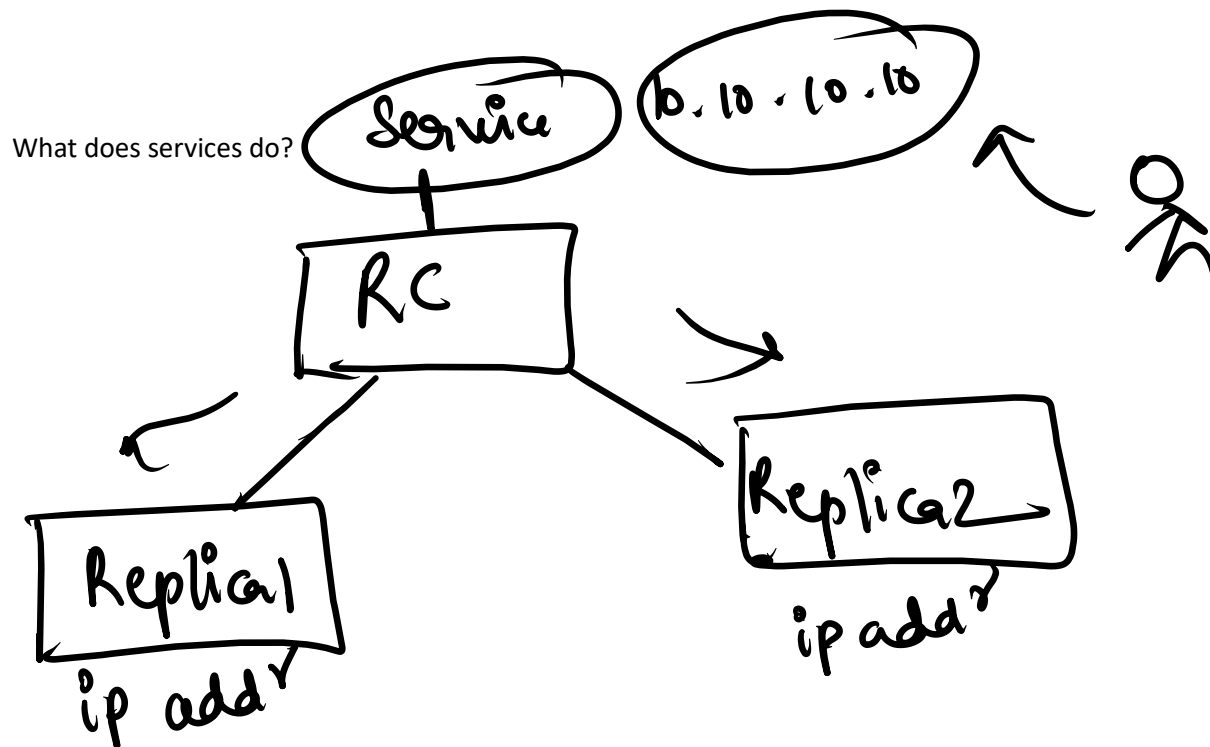
This leads to a problem if some set of pods (call them backend) provides the functionality to other pods (call them frontend) inside your cluster, how do the frontend find out and keep the track of which ip address to connect to , so that the frontend can use the backend part of workload?

FOR THESE WE WILL USE SERVICES

Second reason we should choose services:

We know that pod has a unique ip address , those ip are not exposed outside the cluster

With the service helps to expose the VIP mapped to the pods and allow application to receive traffic from outside.



Type of services:-

- ➔ Cluster ip (default)
- ➔ Node port
- ➔ Load balancer created by cloud providers that will route the external traffic to every node on the nodeport
- ➔ Headless -> creates several endpoints that are used to produce DNS records each DNS record with a pod .

